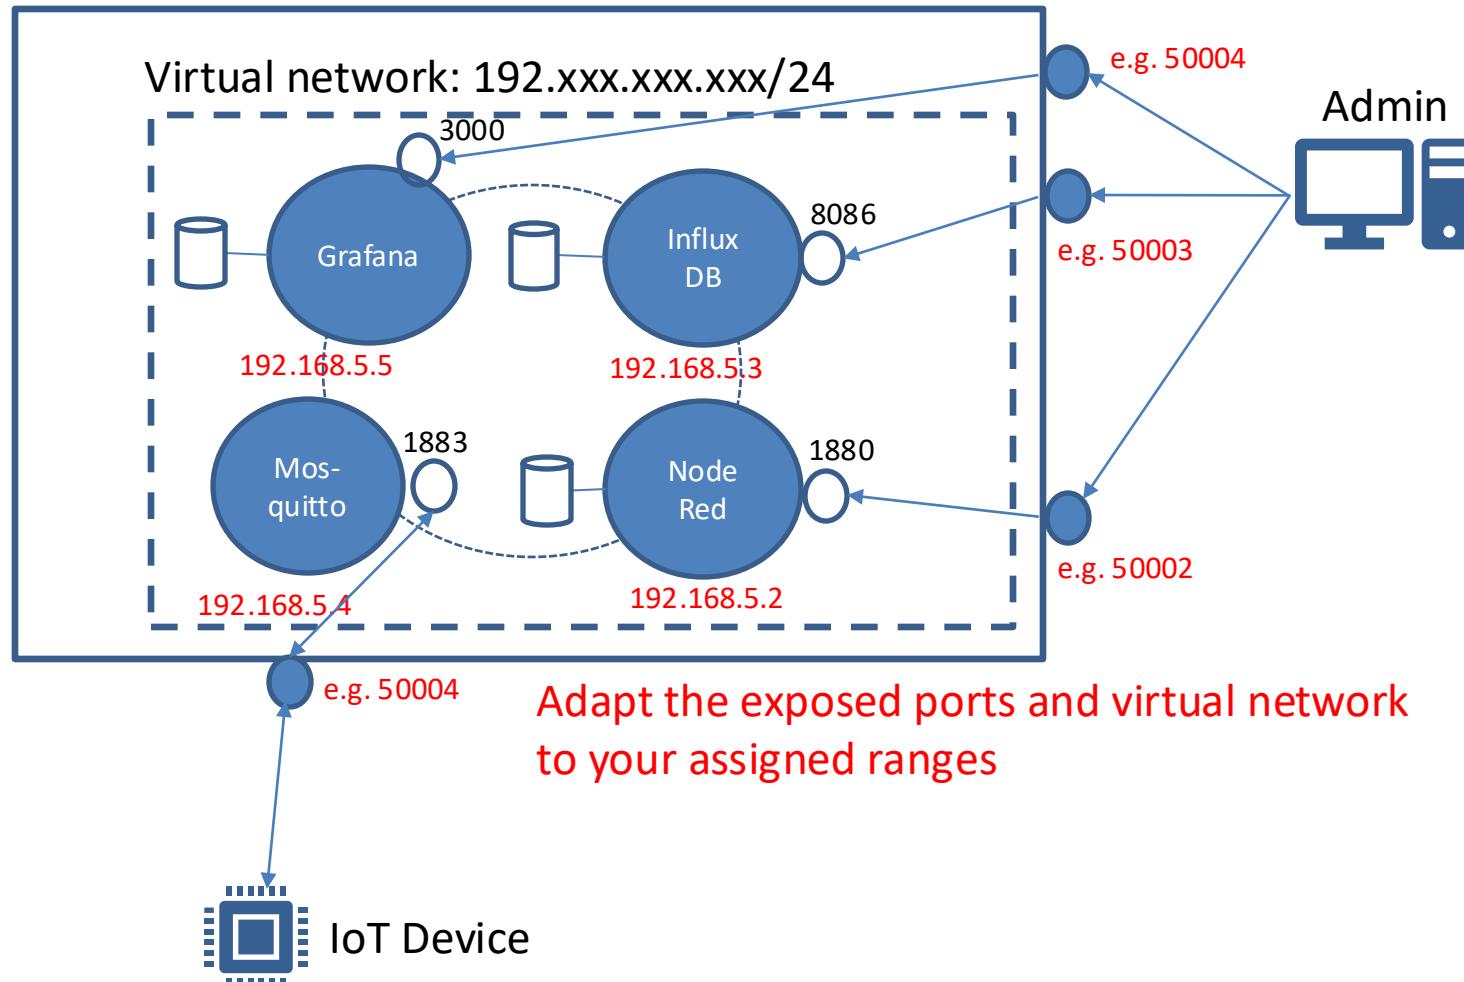


Endpoint on THM EI Cloud: e.g. `iot-lab01.ei.thm.de`



## IoT Lab 2:

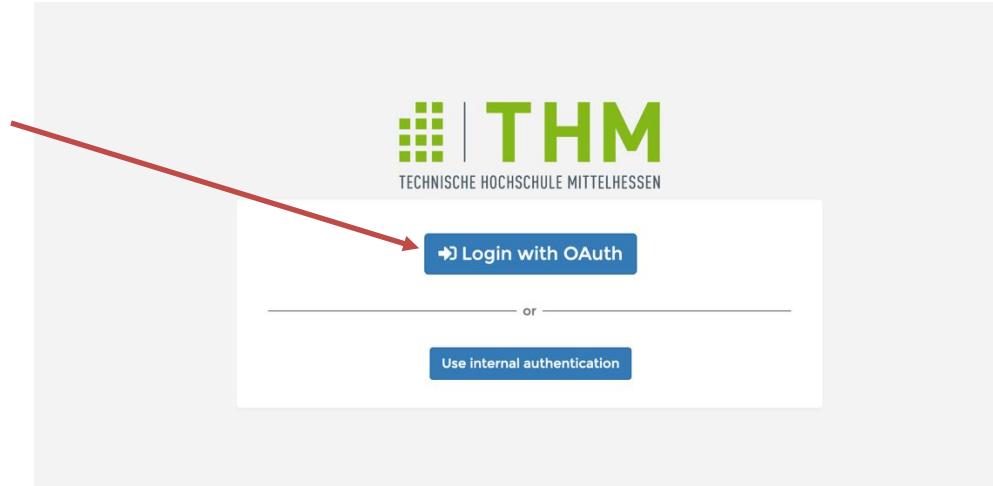
### Deploying Docker Container Apps on the THM EI Cloud

# Learning Objectives

- Using different cloud platform: EI IoT Cloud Platform (portainer)
- Deploy and configure container applications using a **YAML File** (Docker compose):
  - **NodeRed** with a local volume attached for persistent data + secure with password
  - **Mosquitto Broker** (no local volume is necessary)
  - **InfluxDB** with a local volume attached for persistent data
  - **Grafana** with a local volume attached for persistent data
- Use assigned **port-mapping**
- Create an easy „Sensordata“-Flow: publish local CPU load (from lab 1) via MQTT to own broker, subscribe in Node-Red in Cloud, store into database and visualize in grafana
- Subscribe to the defined topics in different ways using NodeRed, MQTT Explorer & MQTTx
- Learn how to secure your MQTT connection via TLS and how to use the secured EI MQTT Broker → used for future labs
- To deploy container there are mutiple options
  - a.) Portainer GUI (standalone or yaml file)
  - b.) Docker Command Line Interface  (*already done in Lab 1*)

# Register and login to THM EI Cloud

- Navigate to <https://cloud.ei.thm.de> \*
- Follow **Login with OAuth**
- Login with SSO with your THM credentials (like in Moodle, etc)
- Your should see an assigned environment, like IoT-Lab0x



- \* Note: If you are not in THM network,  
connect via VPN client

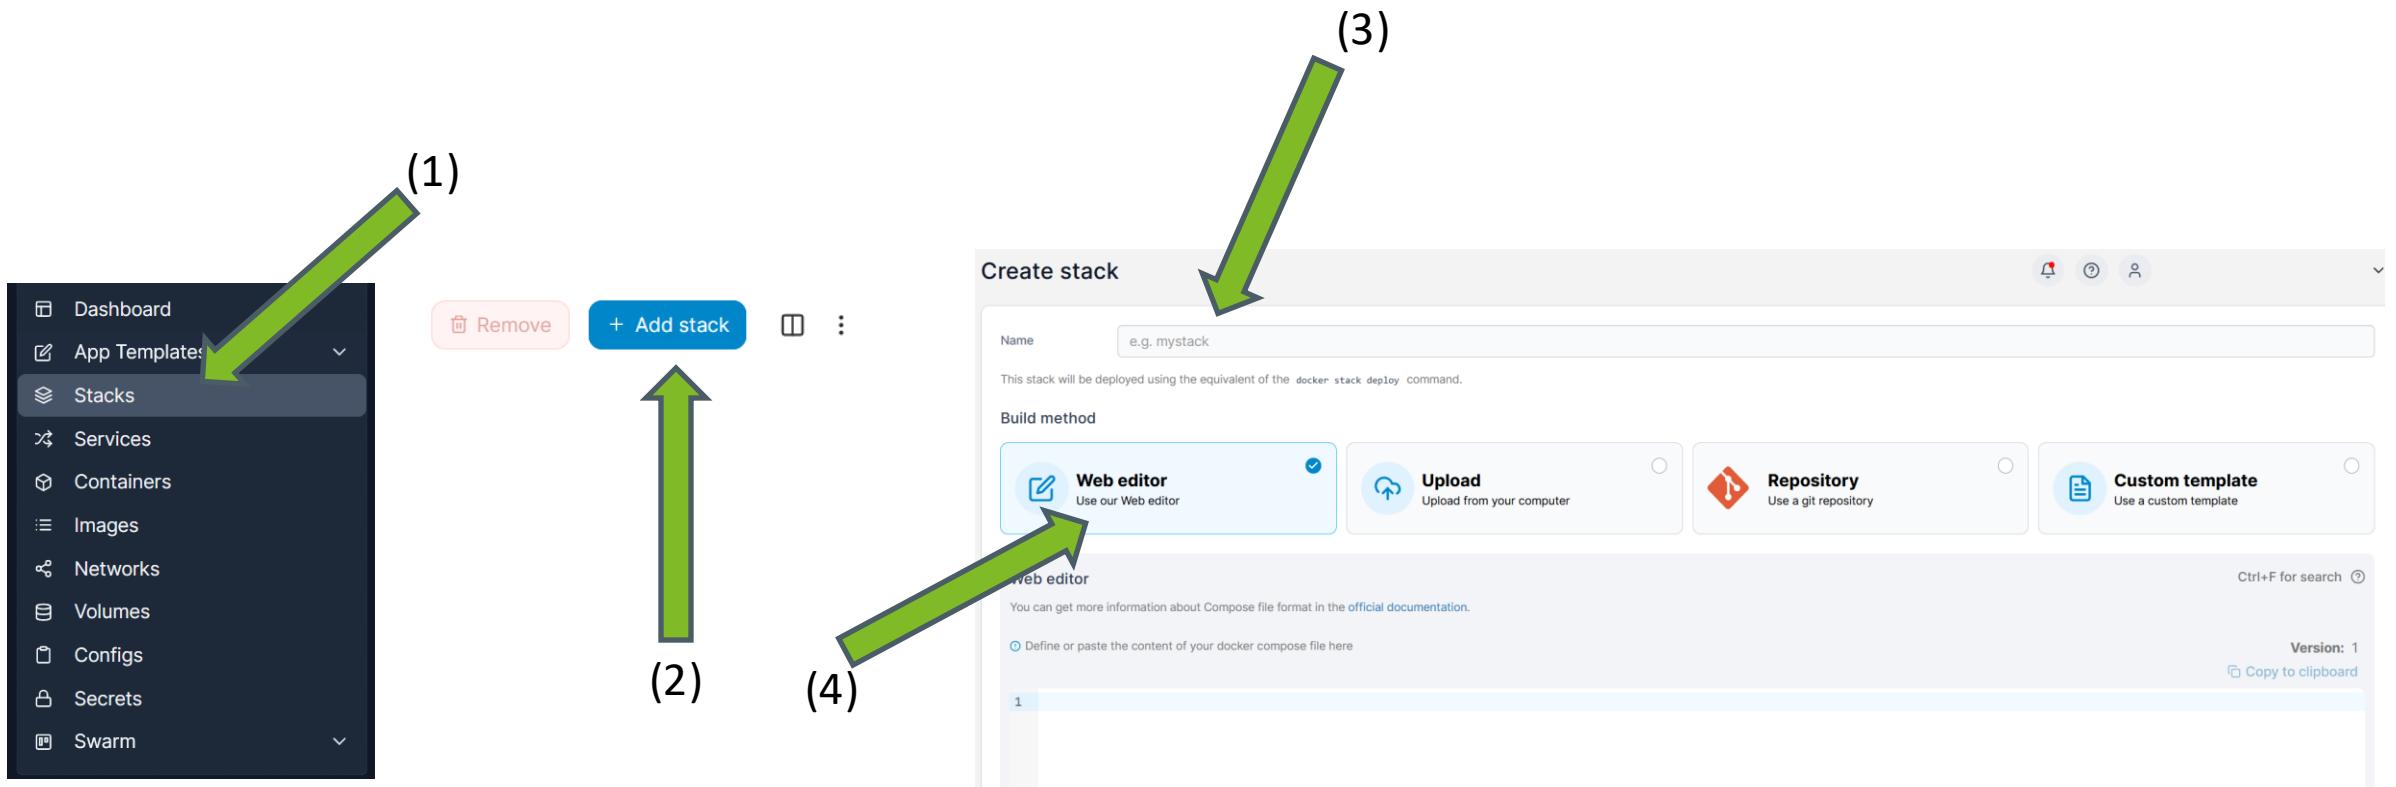
<https://www.thm.de/its/campusnetz/vpn.html>

- Note down your **port-range** assigned to you in moodle / lecture → since you are all working on the same server, ports have to be unique per container!

your port range: \_\_\_\_\_

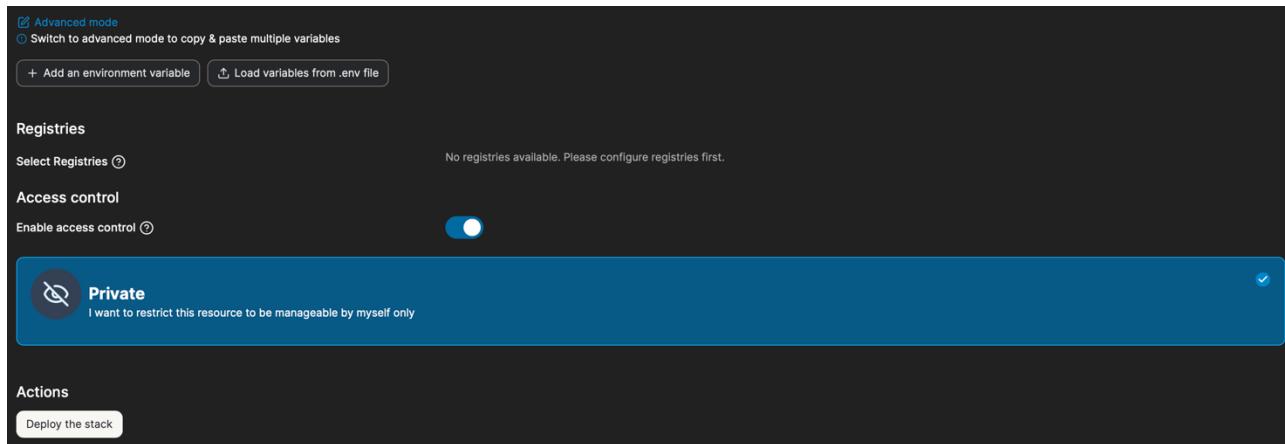
# Create a new stack

- **Login** to your endpoint
  - Head to “Stacks” (1) and click “Add stack” (2)
  - Choose <yourname>\_lab2\_ws20xx as the name of your stack (3)
  - Use the web editor to configure your services (4) → see next slide



# Create a new stack

- Deploy Stack using the docker-compose web editor:
  - Add the services Mosquitto, NodeRed, InfluxDB and Grafana
  - Change <yourname> to your name (lowercase only)
  - Change the yxxxx ports to your own ports assigned by the tutor
  - Ensure that each service has its own unique port; do not reuse ports
  - Check your configuration and deploy the stack  
don't get nervous, it takes about a minute to pull the images ☺
  - DO NOT copy from slides !!! → use the provided example
  - Attention: Indentation spaces are important



```
version: "3.8"
services:
  node-red-<yourname>:
    image: nodered/node-red:latest
    ports:
      - "yxxxx:1880"
    networks:
      - <yourname>_net
    volumes:
      - <yourname>_nodered_vol:/data

  mosquitto-<yourname>:
    image: eclipse-mosquitto:1.6.13
    ports:
      - "yxxxx:1883"
    networks:
      - <yourname>_net

  influxdb-<yourname>:
    image: influxdb:latest
    ports:
      - "yxxxx:8086"
    networks:
      - <yourname>_net
    volumes:
      - <yourname>_influx_vol:/var/lib/influxdb2:rw

  grafana-<yourname>:
    image: grafana/grafana:latest
    ports:
      - "yxxxx:3000"
    networks:
      - <yourname>_net
    volumes:
      - <yourname>_grafana_vol:/var/lib/grafana:rw

volumes:
  <yourname>_nodered_vol:
  <yourname>_influx_vol:
  <yourname>_grafana_vol:

networks:
  <yourname>_net:
```

# YAML File (some theory)

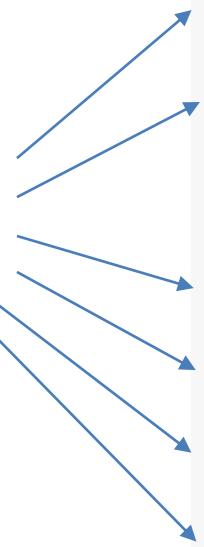
- **YAML** stands for “**YAML Ain’t Markup Language**” — it’s a simple, human-friendly format used to describe data. Think of it like a “recipe” that tells a system what to do.
- In the context of **Docker** or **Kubernetes**, YAML files are used to define:
  - Which containers to run
  - What settings they should have
  - How they should connect to each other (networking)
  - How many copies (scaling)
  - And more...
- **What happens behind the scenes?**
  - Once you write and run the YAML file:
  - Containers are **automatically deployed**
  - Docker (or Kubernetes) **manages** the services
  - **Scaling, networking, restart policies** — all are automated
- **So what do we need to define?**
  - Just **what we want**, like:
    - The container image (e.g., nginx, node, mysql), Port mapping (e.g., 8080:80), Environment variables(opt), volumes, etc.

```
version: "3.9"
services:
  web:
    image: nginx
    ports:
      - "8080:80"
```

# YAML File Format

First level

Remember first level keywords



```
version: "3.7"
services:
networks:
volumes:
configs:
secrets:
```

Web editor

You can get more information about Compose file format in the [official documentation](#).

The code snippet shows a YAML file structure. The first-level keywords are: version, services, networks, volumes, configs, and secrets. Arrows point from the text 'Remember first level keywords' to each of these six keywords.

# YAML File Format

How to write a YAML file:

- No Spaces
- No Upper Case Letters
- Use either Tab key or 2 spaces

Tab key or 2  
spaces



The screenshot shows a web-based YAML editor interface. At the top, there's a blue header bar with the text "Web editor" and "Use our Web editor". Below the header, the title "Web editor" is displayed. A message encourages users to get more information about Compose file format from the "official document". The main area is a code editor containing a YAML configuration file. The file is numbered from 1 to 12. It defines a "services" section with a "nodered" service, which includes an "image" field set to "nodered/node-red:latest" and a "ports" field with a mapping from "6008" to "1880". It also defines a "database" service with an "image" field set to "bitnami/mongodb:latest". Finally, it defines a "networks" section. Vertical lines with arrows point from the numbers 1 through 12 to the corresponding levels of indentation in the YAML code, illustrating the hierarchical structure of the file.

```
1 versions: "3.8"
2
3 services:
4   nodered:
5     image: nodered/node-red:latest
6     ports:
7       - 6008:1880
8   database:
9     image: bitnami/mongodb:latest
10
11 networks:
12
```

# Access your applications

- Note down your IPs / Ports to access your applications:

- Node-Red

Public IP: \_\_\_\_\_  
Private IP: \_\_\_\_\_

Public Port: \_\_\_\_\_  
Private Port: \_\_\_\_\_

- Mosquitto

Public IP: \_\_\_\_\_  
Private IP: \_\_\_\_\_

Public Port: \_\_\_\_\_  
Private Port: \_\_\_\_\_

- InfluxDB

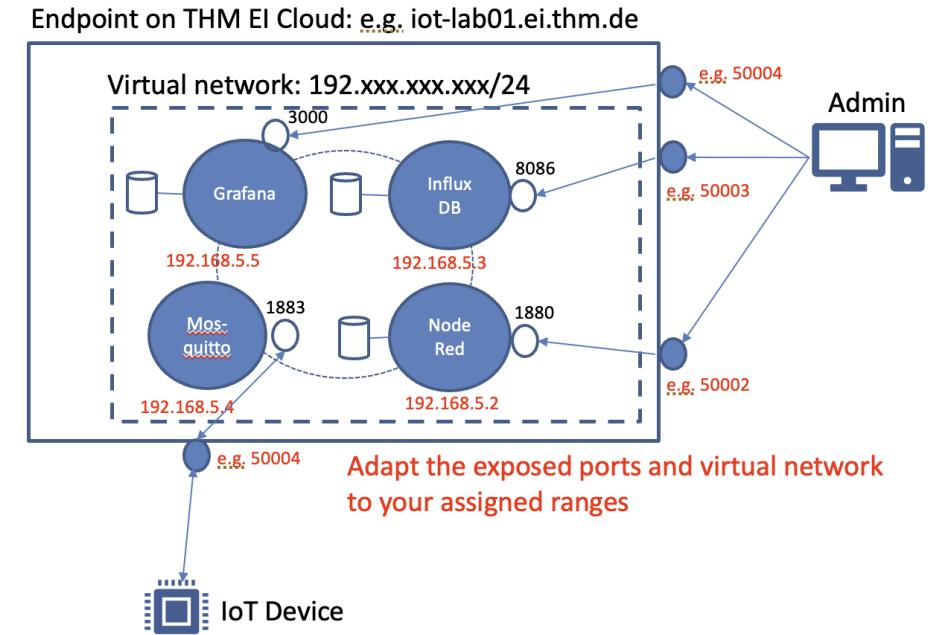
Public IP: \_\_\_\_\_  
Private IP: \_\_\_\_\_

Public Port: \_\_\_\_\_  
Private Port: \_\_\_\_\_

- Grafana

Public IP: \_\_\_\_\_  
Private IP: \_\_\_\_\_

Public Port: \_\_\_\_\_  
Private Port: \_\_\_\_\_

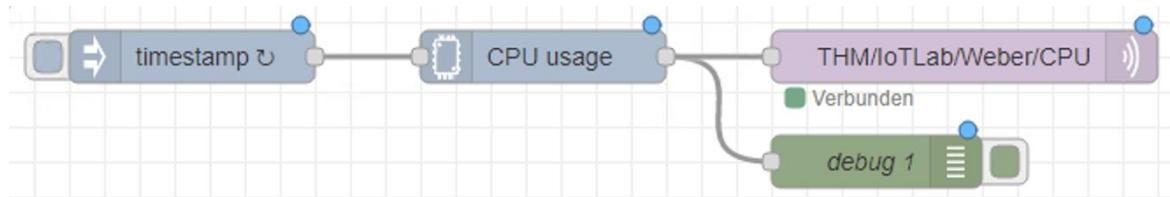


# Configure Node Red & MQTT

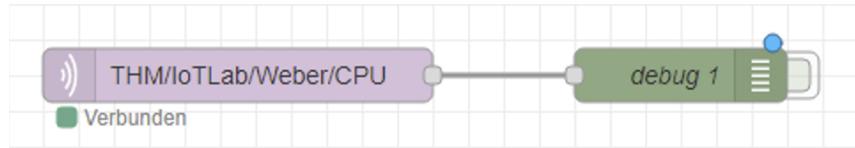
- Open **NodeRed** on your local computer (from lab 1)
  - Publish the CPU Data to either your broker (a) or EI Broker (b)
  - (a) use your own broker via

Public IP: \_\_\_\_\_ Public Port: \_\_\_\_\_  
Private IP: \_\_\_\_\_ Private Port: \_\_\_\_\_

(b) Use EI Broker: **mqtt.ei.thm.de:1993** user/pw: **iotlab/iotlab** Topic format: **THM/IoTLab/yourname**



- Open **NodeRed deployed on cloud (this lab)**
  - subscribe to your published data and show values in debug

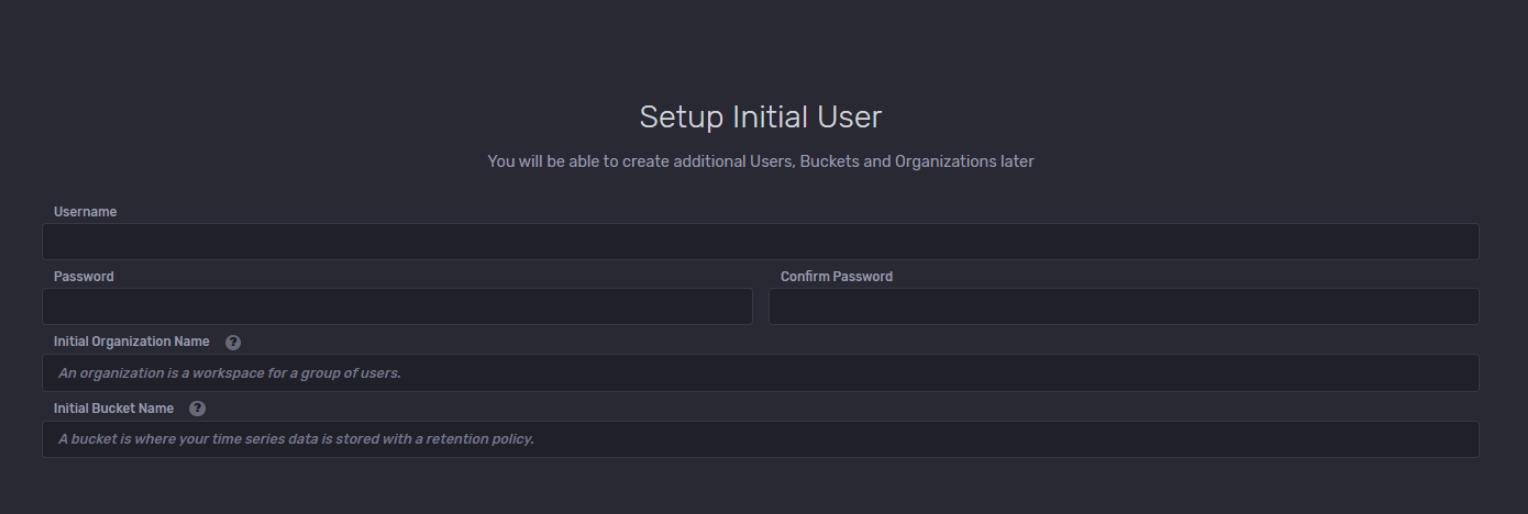


```
9.8.2022, 14:51:58 node: debug 1
THM/IoTLab/Weber/CPU : msg : Object
object
topic: "THM/IoTLab/Weber/CPU"
payload: 0.5833333333333334
qos: 2
retain: false
msgid: "0555fba9bacdf2ef"
```

# Configure InfluxDB

## ■ Open InfluxDB

- Walk through the Get Started guide and setup initial user
  - Choose a username and password and **note it down**
  - Choose the Organization Name: THM
  - Name your data bucket: lab02



Setup Initial User

You will be able to create additional Users, Buckets and Organizations later

Username

Password

Confirm Password

Initial Organization Name ?

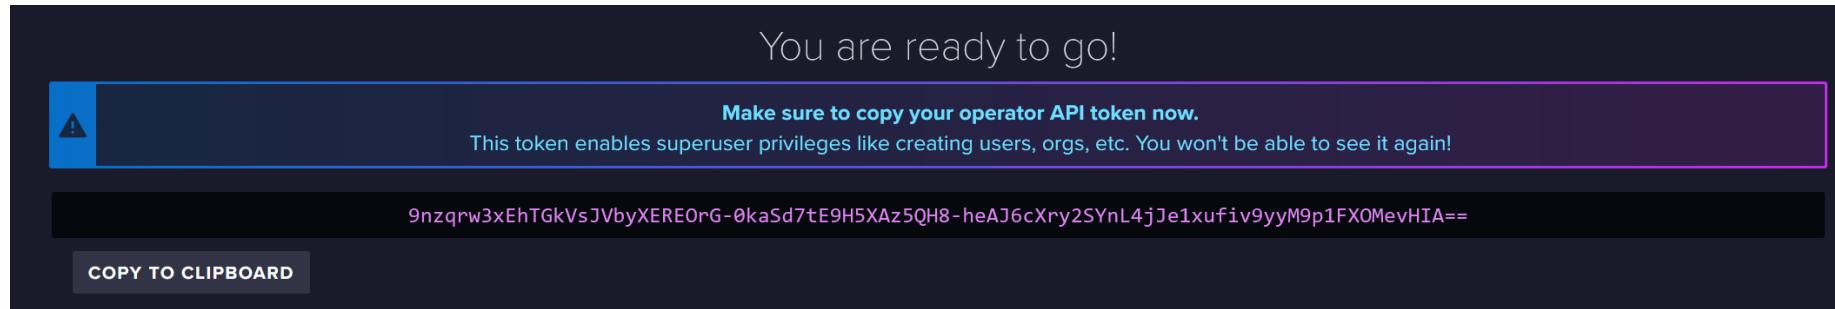
An organization is a workspace for a group of users.

Initial Bucket Name ?

A bucket is where your time series data is stored with a retention policy.

# Configure InfluxDB

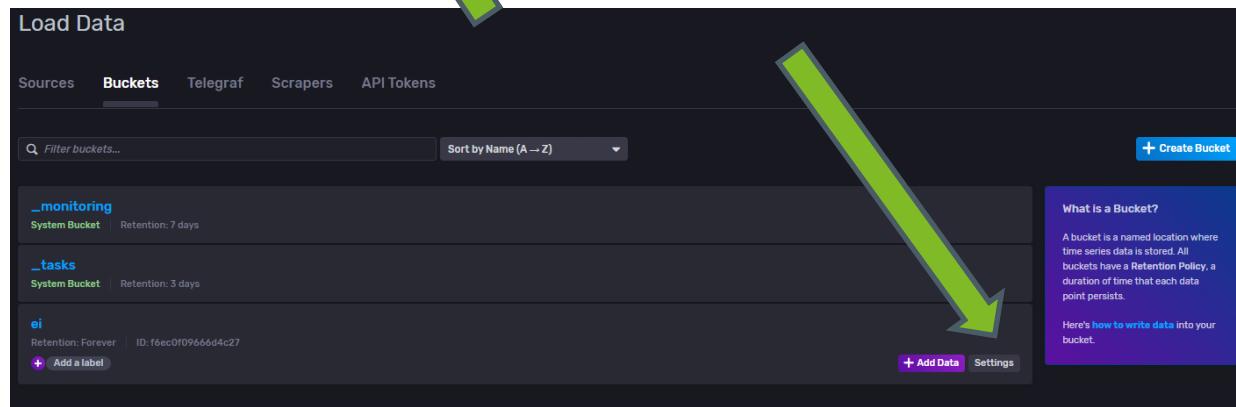
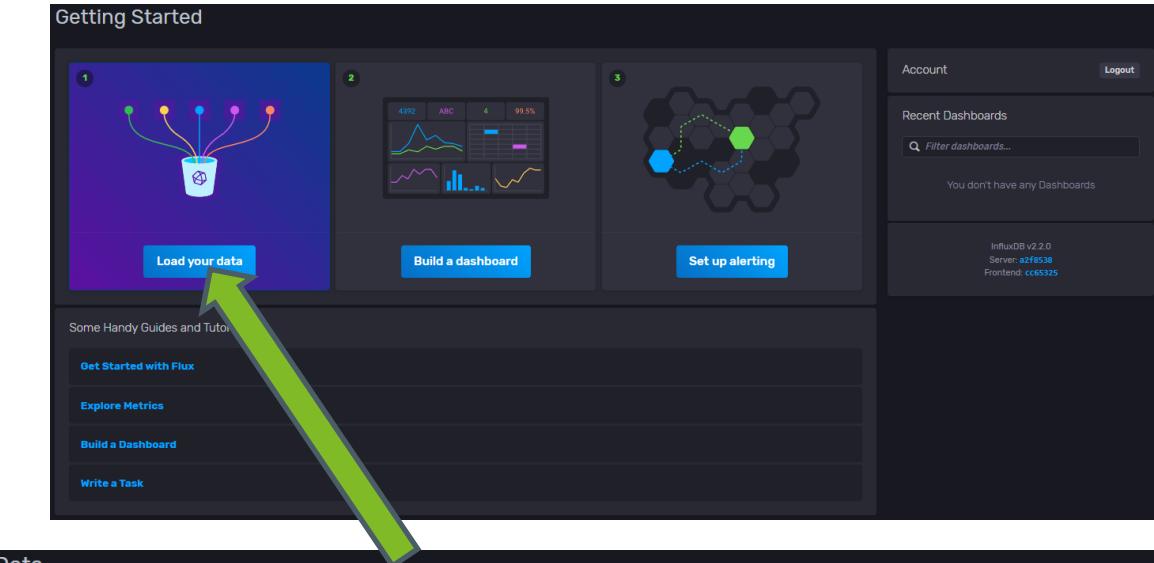
- **Save your InfluxDB API key (use e.g. editor)**
  - Make a note of the API key, as we'll be using it later.
  - → COPY TO CLIPBOARD not always working, so copy manually !!



# Configure InfluxDB

## Configure InfluxDB

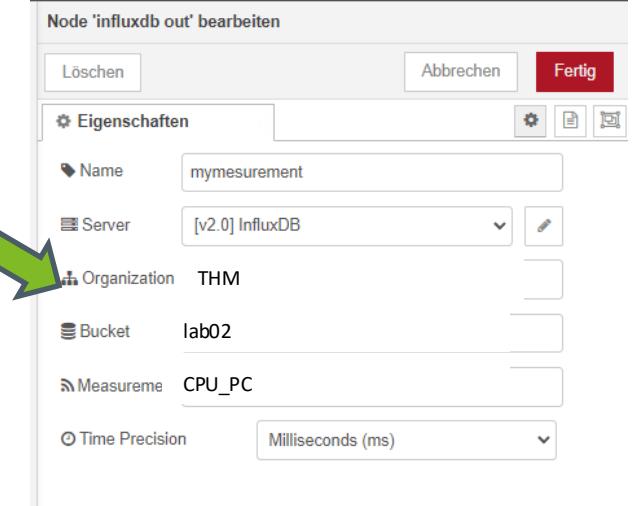
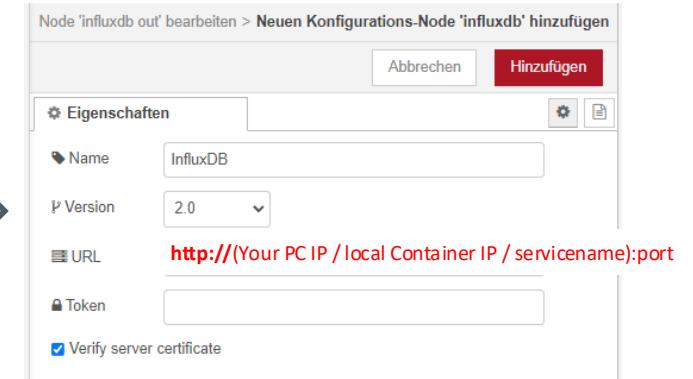
- Click “Load your data” on the Getting Started Page
- In the “Data” menu, head to Buckets and change the retention policy in “settings” if needed



# Configure InfluxDB in NodeRed

## Add InfluxDB Server to Node Red (cloud)

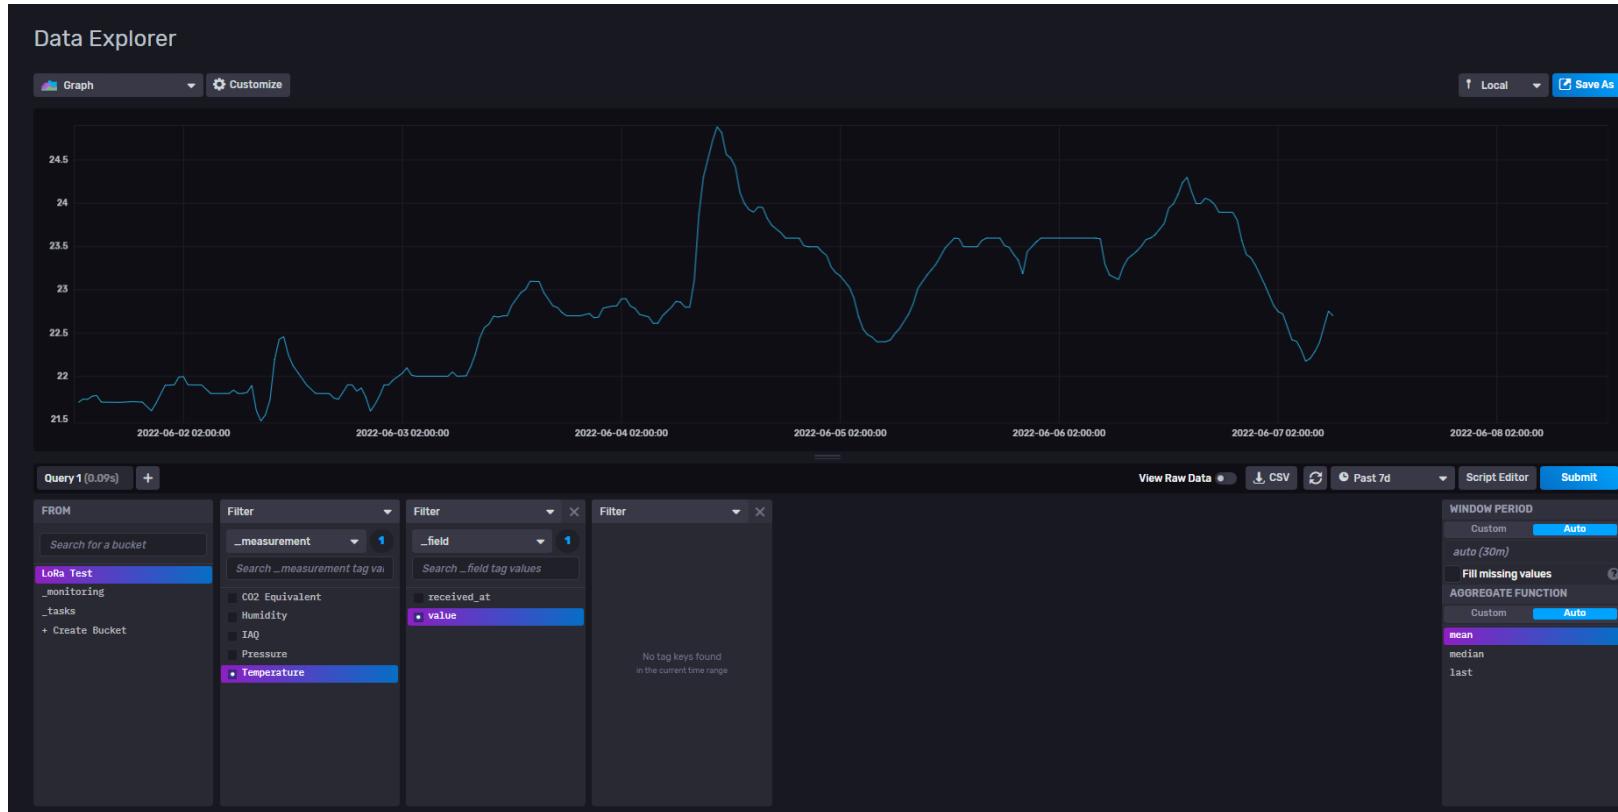
- Go back to your Node Red Service and add an influxdb out node
  - needs to be installed → manage palette
- Configure a new InfluxDB Server and select version 2
- Enter your URL (http:// + IP/FQDN cloud endpoint or local Container IP or service name + port of InfluxDB) and the copied admin token
- Add the configuration
- Edit the influxdb out node and enter your organization, bucket, measurement
  - Organization: THM
  - Bucket: lab02
  - Measurement: CPU\_PC



# Configure InfluxDB

## View your measurements

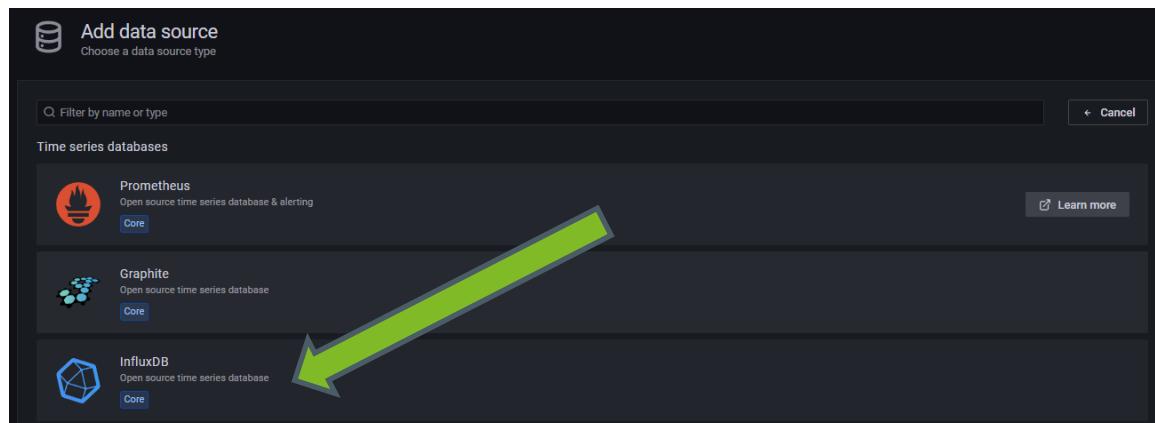
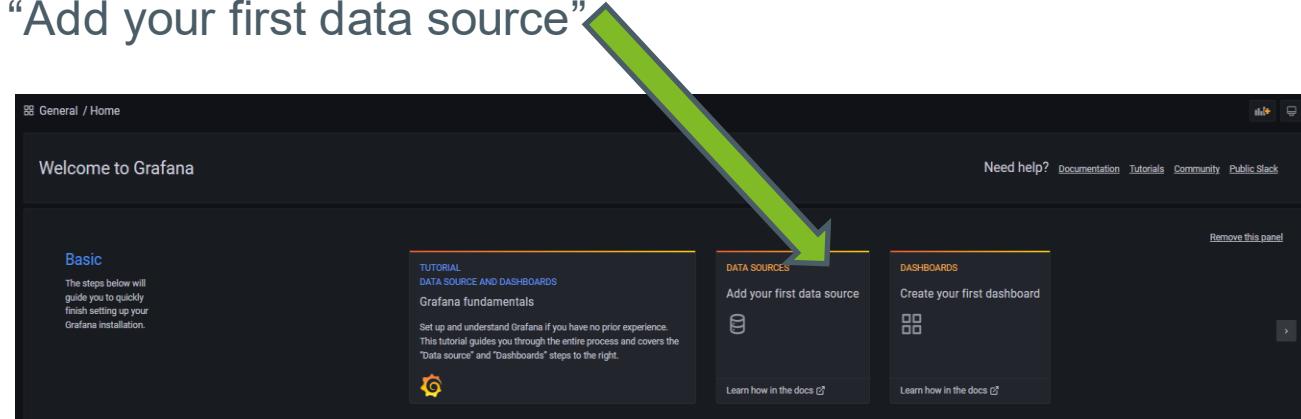
- On your InfluxDB Web-UI go to the Explore menu
- You can now select your bucket and view your measurements and click submit



# Configure Grafana

## ■ Open the Grafana Web Interface

- Login with the default user/password combination (admin/admin)
- Choose a new password and **note it down**
- On the Grafana Homepage click the “Add your first data source” Panel
- Add InfluxDB as a source



# Configure Grafana

- Add InfluxDB to Grafana
  - Change the Query Language to **Flux**
  - Enter the URL and Port of your InfluxDB
  - Disable Basic Auth
  - Enter your **Organization, Token and Default Bucket** under "InfluxDB Details"
  - Click save and test → should be green and successful

InfluxDB Details

Organization	THM
Token	configured
Default Bucket	lab02
Min time interval	10s
Max series	1000

**Reset**

HTTP

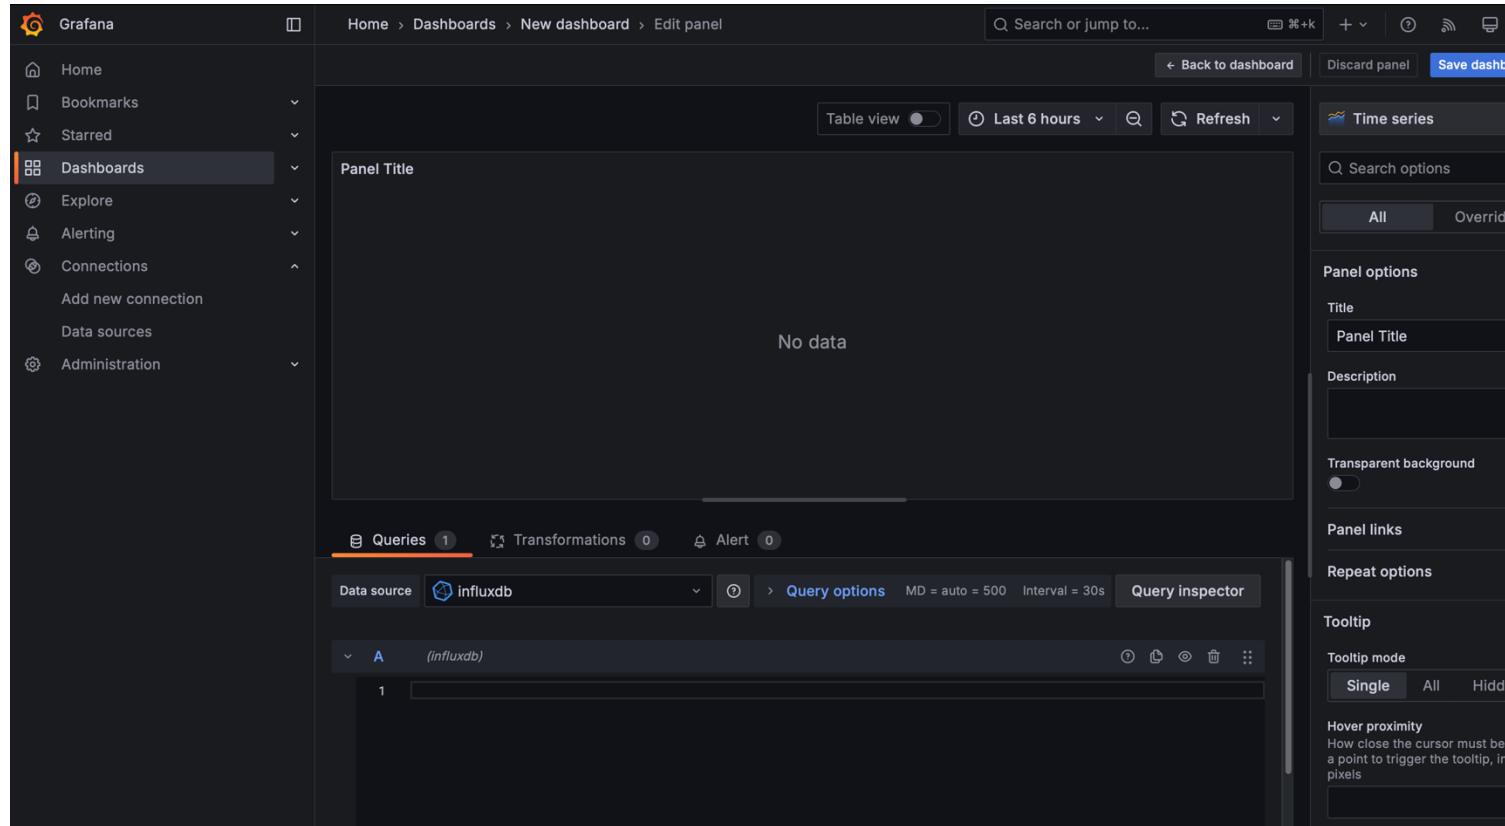
URL	Your IP and port
Allowed cookies	New tag (enter key to add) <b>Add</b>
Timeout	Timeout in seconds

Auth

Basic auth	<input checked="" type="checkbox"/> With Credentials <input type="checkbox"/>
TLS Client Auth	<input checked="" type="checkbox"/> With CA Cert <input type="checkbox"/>

# Configure Grafana

- Create a dashboard and display measurements
  - Click on the “Create your first dashboard” panel at the Homepage of Grafana
  - Add visualization and select your influxdb as data source



# Configure Grafana

- Go back to your InfluxDB Explorer
  - Choose your bucket and a measurement you want to visualize (value)
  - Click on the script editor button
  - Copy the query text for Grafana

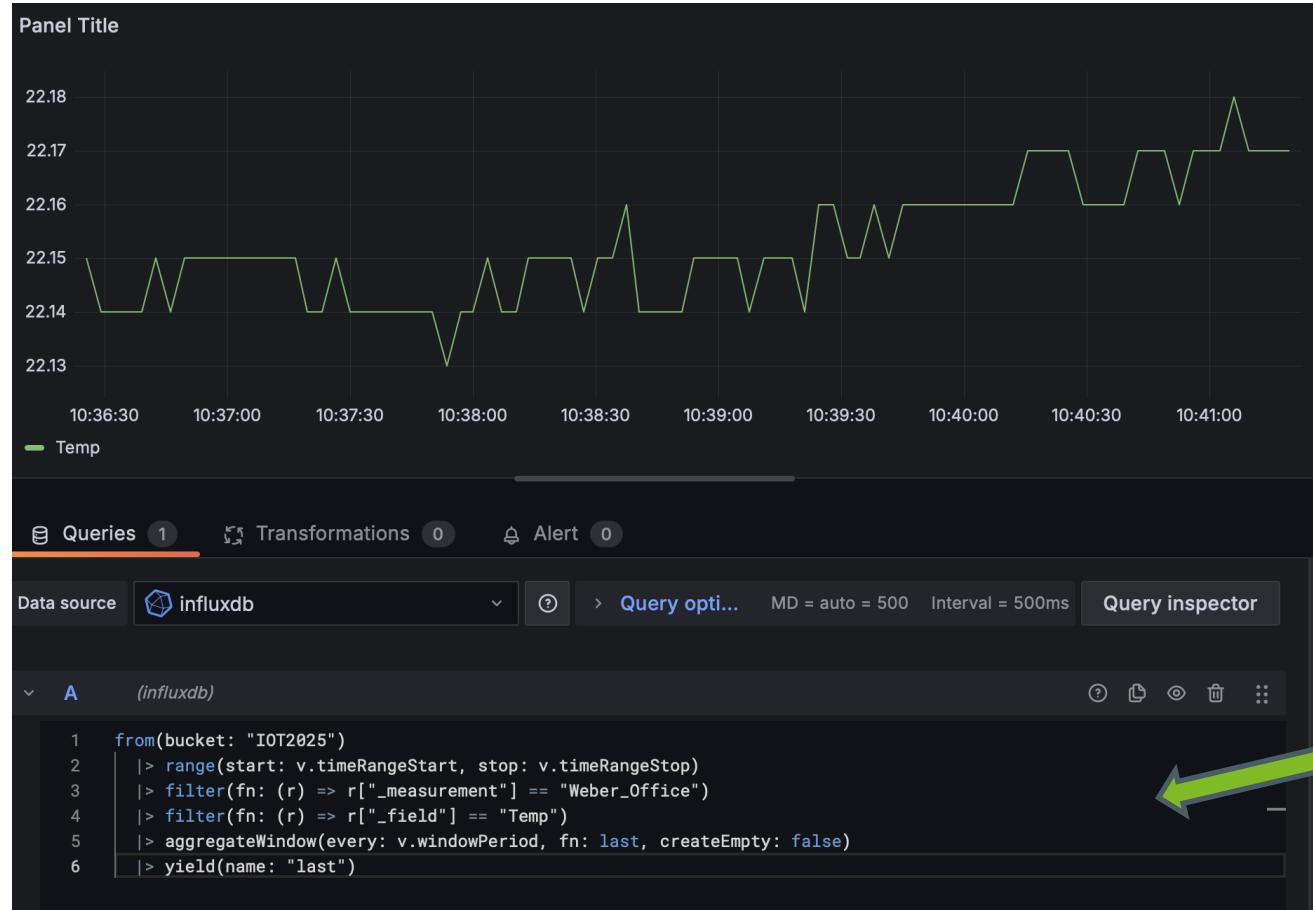
The screenshot shows the Grafana interface with a line chart visualization. The chart displays a single data series over time, with values ranging from 22.11 to 22.14. The x-axis shows timestamps from 2025-03-11 10:20:00 to 2025-03-11 10:30:00. Below the chart, the InfluxDB query editor is visible. The query text is:

```
1 from(bucket: "IOT2025")
2 |> range(start: v.timeRangeStart, stop: v.timeRangeStop)
3 |> filter(fn: (r) => r["_measurement"] == "Weber_Office")
4 |> filter(fn: (r) => r["_field"] == "Temp")
5 |> aggregateWindow(every: v.windowPeriod, fn: last, createEmpty: false)
6 |> yield(name: "last")
```

A large green arrow points from the text "Copy the query text for Grafana" in the list above to the "SCRIPT EDITOR" button in the Grafana interface.

# Configure Grafana

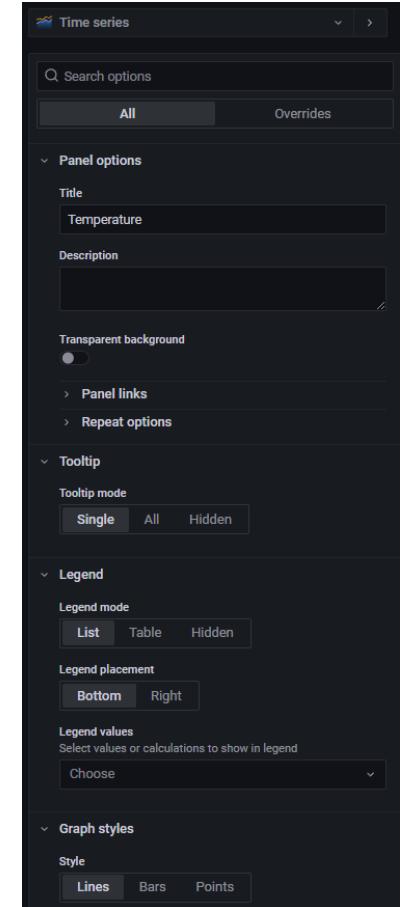
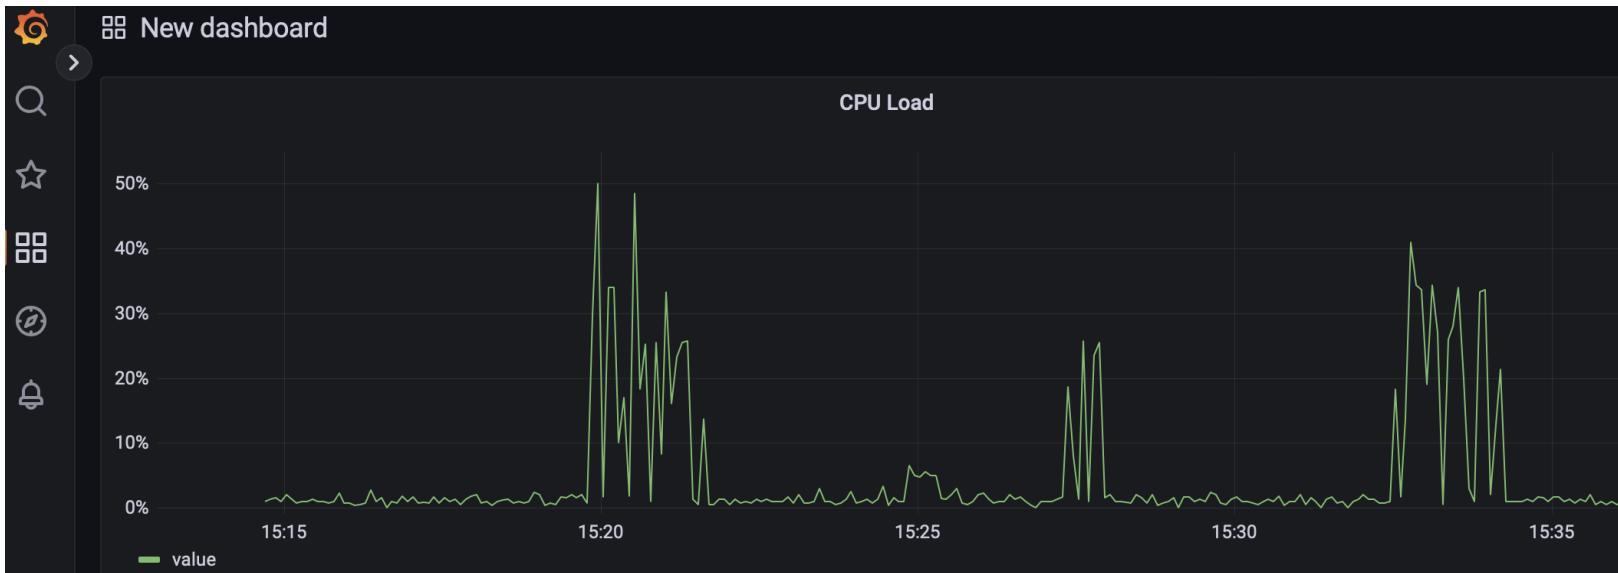
- Insert the copied Query text in your Grafana raw query editor



```
1 from(bucket: "IOT2025")
2 |> range(start: v.timeRangeStart, stop: v.timeRangeStop)
3 |> filter(fn: (r) => r["_measurement"] == "Weber_Office")
4 |> filter(fn: (r) => r["_field"] == "Temp")
5 |> aggregateWindow(every: v.windowPeriod, fn: last, createEmpty: false)
6 |> yield(name: "last")
```

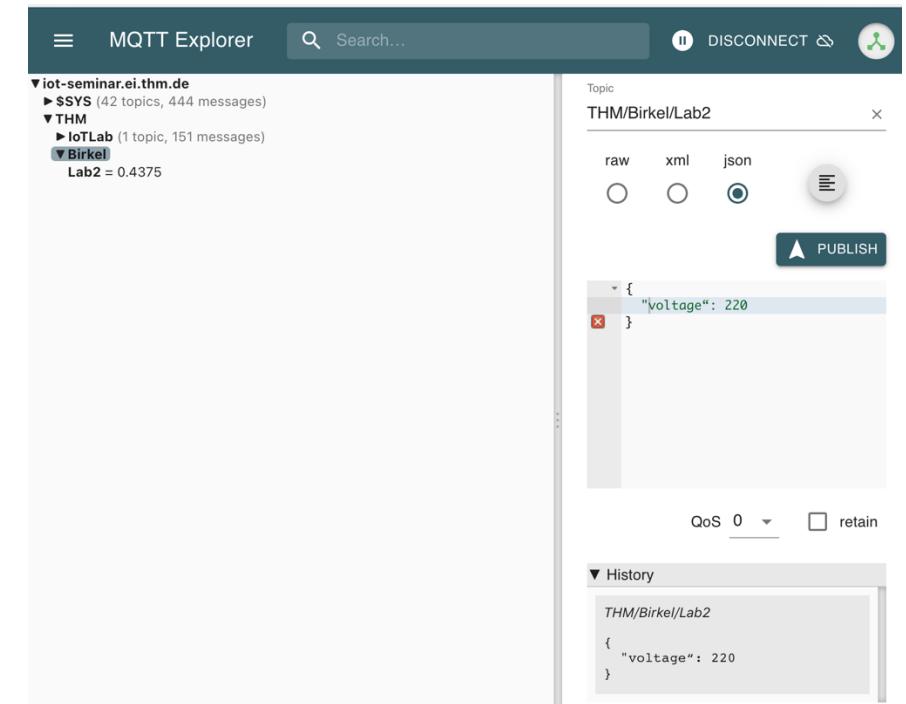
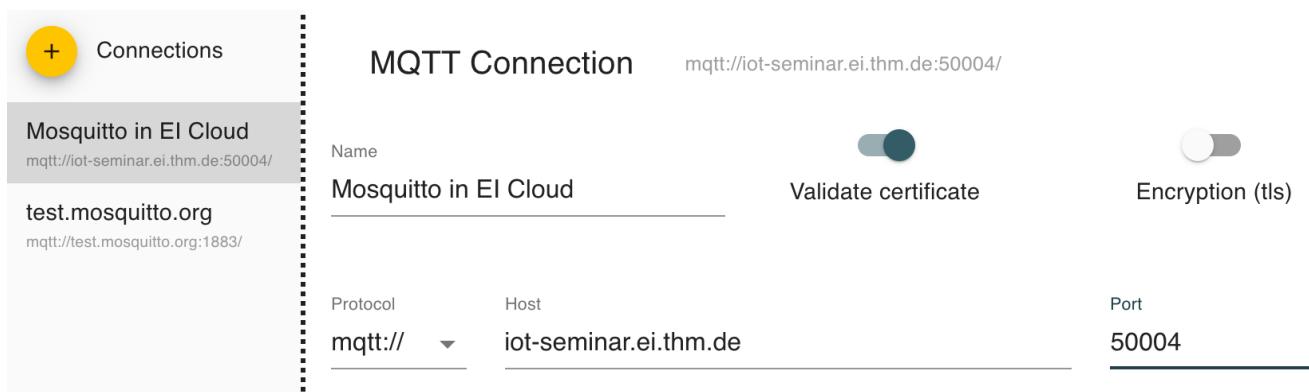
# Configure Grafana

- Now your turn:
  - Give a title, description and the correct unit
    - Fine tune your graph if you want to
    - Click apply if you are done and **save your dashboard**  
It will not be saved automatically!



# Experiment with MQTT Clients

- **MQTT Explorer**
- **Open MQTT Explorer (already installed in Lab)**
  - Add a new connection Server: either your own broker or EI broker
  - Subscribe to your topic and show your data

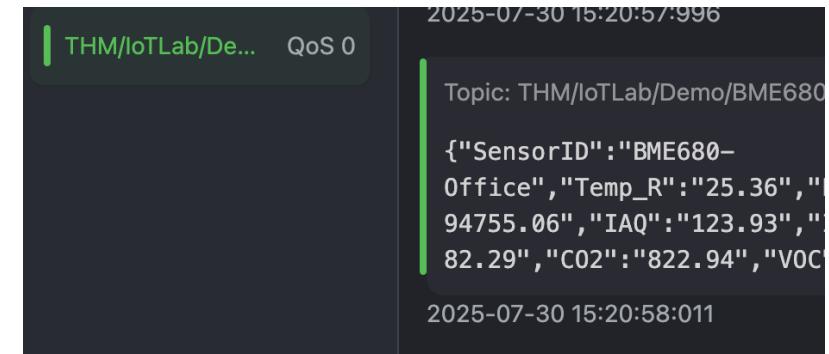


# Experiment with MQTT Clients

- MQTTx
- Open MQTTx (now we use TLS on EI Broker)
  - Add a new connection Server: either your own broker or EI broker
  - Subscribe to your topic and show your data

The screenshot shows the MQTTx configuration interface. The connection details are as follows:

- Name: EI\_Broker (student)
- Host: mqtts:// mqtt.ei.thm.de
- Port: 9993
- Client ID: mqttx\_b3846970
- Username: iotlab
- Password: ..... (redacted)
- SSL/TLS: Enabled
- SSL Secure: Enabled
- ALPN:
- Certificate: CA signed server certificate (selected)



# Add Authentication to NodeRed

- **Option 1: Modify setting.js file to define username&password**
- Go to your NodeRed volume and download setting.js, go to // securing node red section in this file
- Uncomment this string including the users key value which has one item defined as a JSON Object
  - JSON Objects defines: Username, password and permissions
  - We can generate the hash of the password via a hashfunction:  
`password: require('bcryptjs').hashSync("cloud2021")`  
Libary bccrypt includes the hashSync function which returns the hash value of the password
  - Rename old setting.js to setting.js.backup and upload modified file including the password
  - Go to your stack, choose NodeRed and update it => NodeRed will be restartet and reads the settings.js file
  - Login with your username and password
  - **Don't forget to logout in NodeRed on the upper right corner!**



- **Not used now: Option 2: Using an environment Variables to define username & pw**
- Environment variables exist at the runtime of the container
- Modify your deployment stack file (YAML) and define username and password
- Update your settings.js file to read the new environment variable:
  - Go to your nodered volume
  - Download settings.js file
  - Rename the original setting.js file into setting.js.backup
  - Modify it according to the slide on the next page & upload it again
  - Restart the NodeRed app and login with the username and password to node red
  - Modify the password in the deployment file and test if it works!

```
// Securing Node-RED
// -----
adminAuth: {
  type: "credentials",
  users: [
    {
      username: "admin",
      password: require('bcryptjs').hashSync("cloud2021"),
      permissions: "*"
    }
  ],
}
```

```
version: "3.8"

services:
nodered1:
  image: nodered/node-red:latest
  ports:
  - "50002:1880"
environment:
  - NODERED_MY_PASSWORD=cloud2021
  - NODERED_MY_USERNAME=me
networks:
  network2:
volumes:
  - nodered1-v:/data
```

# Modifying the file settings.js on volume

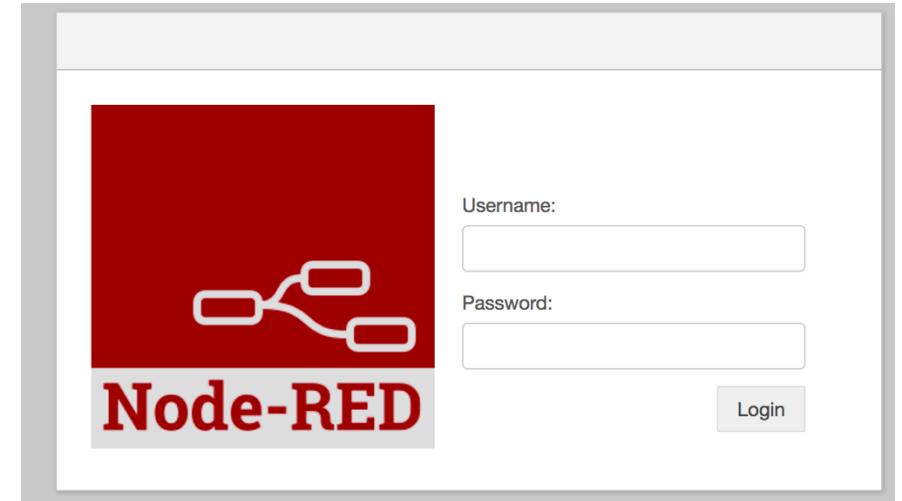
//All the way on the to in settings.js enter the following function:

```
my_function = ()=> {

    username = process.env.NODERED_MY_USERNAME
    password = process.env.NODERED_MY_PASSWORD

    return {
        username,
        password
    }
}

// Securing Node-RED
// -----
// To password protect the Node-RED editor and admin API, the following
// property can be used. See http://nodered.org/docs/security.html for details.
adminAuth: {
    type: "credentials",
    users: [
        {
            username: my_function()['username'],
            password: require('bcryptjs').hashSync(my_function()['password']),
            permissions: "*"
        }
    ],
}
```



# Lab Summary

## EI IoT Cloud platform, Docker compose, secure IoT services and MQTTS

In this lab, you learned to:

### ■ Deploy IoT Applications in the Cloud

- Use Portainer GUI and Docker Compose YAML for container deployment
- Manage Node-RED, Mosquitto, InfluxDB, and Grafana with volume and port configuration + securing NodeRed with username and password

### ■ Build and Visualize IoT Data Flows

- Publish system metrics via MQTT
- Subscribe in Node-RED (cloud)
- Store in InfluxDB
- Visualize using Grafana dashboards

### ■ Work Securely with MQTT

- Subscribe using multiple clients (Node-RED, MQTT Explorer, MQTTx)
- Secure connections with TLS
- Use the **EI MQTT Broker** with certificates



You now understand

- multi-container deployment
- cloud workflows
- MQTT and service security
- → all essential for scalable IoT solutions.

# Lab 2 Review & Reflection

- What is the purpose of Docker Compose?
- Which services did you deploy using a YAML file?
- Why is port mapping necessary in container-based applications?
- How do you persist data for containers like Node-RED, InfluxDB, and Grafana?
- What is the purpose of an MQTT Broker? Which are the well-known ports?

# Lab 2 Review & Reflection

- What is TLS and why is it important for MQTT communication?
- Which elements are used to secure connection to the used MQTT Broker?
- Create a topology for this lab (template is provided)